



VIT[®]

BHOPAL

AIML Project Report

Name: Arun Chaudhary

Registration Number: 25BCE10971

Course: CSA2001 Fundamentals in AIML

Program: B. Tech Computer Science and Engineering

Submitted To: Dr. MK Jayanthi

Project Report:

Predicting Startup Profitability with Machine Learning

1. What's This Project About?

The goal of this project was to figure out what really drives a startup's financial success. To do this, I built a Multiple Linear Regression model that looks at how different startups spend their money—specifically on Research & Development (R&D), Administration, and Marketing—and where they are located. By feeding this data into the model, we can accurately predict a startup's revenue growth and understand which investments actually pay off.

2. The Data We Used

I worked with a dataset containing the financial records of 50 different startups. Fortunately, the data was super clean with no missing values.

Here is what we tracked for each company:

- **R&D Spend:** How much money went into researching and developing their product.
- **Administration:** Their basic day-to-day operational costs.
- **Marketing Spend:** Their advertising and promotional budget.
- **State:** Where they operate (California, Florida, or New York).
- **Profit:** The actual net profit they made (this is what our model is trying to predict).

3. What Did the Raw Data Tell Us?

Before even training the model, just looking at the data revealed some pretty clear patterns:

- **The Big Winner:** The most profitable startup in the dataset is based in New York, bringing in ₹192,261.83. Unsurprisingly, they also spent the most on R&D (₹165,349.20) and Marketing (₹471,784.10).
- **The Lowest Performer:** On the flip side, the least profitable startup (based in California) made only ₹14,681.40. The catch? They invested absolutely ₹0.00 in R&D, putting all their money into Admin and Marketing instead.
- **Admin Costs Don't Change Much:** Whether a startup was at the top of the list or the bottom, their administration costs stayed surprisingly flat (mostly between ₹116k and

₹136k). This tells us that admin is just a baseline cost of doing business, not something that drives profit.

4. How I Built the Model

To get the data ready for the machine learning algorithm, I had to make a few tweaks:

1. **Handling the Locations:** Since machine learning models need numbers, not words, I used one-hot encoding to turn the "State" column into numerical data. I dropped California from the variables to avoid a common statistical issue called the "dummy variable trap."
2. **Splitting the Data:** I took 80% of the startups (40 of them) to teach the model how to make predictions. I hid the remaining 20% (10 startups) to test the model later and see how well it actually learned.
3. **Training:** I used Scikit-Learn's Multiple Linear Regression algorithm to find the mathematical relationships between all these variables.

5. Did the Model Work?

Yes, it performed really well. I measured its accuracy using the R^2 score and Root Mean Squared Error (RMSE).

- **Training Accuracy (R^2):** 95.37%
- **Testing Accuracy (R^2):** 89.87%

Basically, the model can explain about 90% of the variance in the unseen test data. The fact that the training and testing scores are so close tells us that the model isn't just memorizing the data (overfitting)—it genuinely learned the underlying patterns and can generalize well to new startups.

6. The Real Drivers of Profit

Once the model was trained, I looked at the "weights" (coefficients) it assigned to each feature. This is where it gets interesting:

- **R&D is King:** For every ₹1 a startup spends on Research & Development, their profit increases by about ₹0.80. It is by far the biggest driver of success.
- **Marketing Helps a Little:** Every ₹1 spent on Marketing adds about ₹0.03 to the profit. It's a positive impact, but much smaller than R&D.
- **Admin Hurts Slightly:** There is a tiny negative correlation with administration costs, meaning bloated operational spending slightly eats into profits.
- **Location Doesn't Matter Much:** Being in Florida or New York gave a tiny mathematical bump to the profit compared to California, but honestly, the geographic impact is practically zero compared to how a startup allocates its budget.

7. Taking It for a Spin

To make sure the model makes logical sense in the real world, I fed it a hypothetical startup with the following profile:

- R&D Spend: ₹100,000
- Admin Spend: ₹120,000
- Marketing Spend: ₹300,000
- Location: New York

The model predicted this startup would generate ₹**135,300.12** in profit, which aligns perfectly with the trends we saw in the original dataset.

8. Final Thoughts

The biggest takeaway from this project is that **building a better product pays off the most**. Startups that aggressively fund their Research & Development consistently outperform those that don't, regardless of what state they operate in. The model I built is statistically solid and serves as a highly reliable tool for forecasting a venture's future financial success based on its budget.